

Wilka Torrico De Carvalho, *Cognitive Scientist and Reinforcement Learning Researcher*

CONTACT INFORMATION	Website: cogscikid.com	Github: github.com/wcarvalho
	E-mail: wcarvalho@g.harvard.edu	Google Scholar
EMPLOYMENT	Harvard , Boston, MA Research Fellow	<i>Sept 2023 – Present</i>
	DeepMind , London, UK Research Scientist Intern	<i>Sept 2022 – May 2023</i>
	DeepMind , London, UK Research Scientist Intern	<i>Mar 2021 – July 2021</i>
	Microsoft Research , Redmond, WA Research Scientist Intern	<i>June 2018 – Aug 2018</i>
	IBM Research , San Jose, CA Research Scientist Intern	<i>Sept 2017 – Dec 2017</i>
	VISA Research , Palo Alto, CA Research Scientist Intern	<i>May 2017 – Aug 2017</i>
EDUCATION	University of Michigan–Ann Arbor , Ann Arbor, MI <i>School of Engineering</i> , Ph.D. in Computer Science Advisors: Satinder Singh , Honglak Lee , Richard Lewis	Sep 2018 - July 2023
	University of Southern California , Los Angeles, CA <i>Viterbi School of Engineering</i> , M.S. in Computer Science Advisor: Yan Liu	Aug 2015 - May 2017
	Stony Brook University , Stony Brook, NY <i>College of Arts and Sciences</i> , B.S. in Physics Advisor: Axel Drees	Aug 2011 - May 2015
	Brooklyn Technical High School , Brooklyn, NY Diploma in Applied Physics	May 2007 - May 2011
HONORS & AWARDS	Kempner Institute Neuro-AI Fellowship	2023
	ICML, Neurips Top Reviewer	2022, 2023
	University of Michigan “The Glass is Half-Full” Award	2019
	1/200 chosen internationally for Heidelberg Laureate Forum	2018
	GEM National Fellowship sponsored by IBM, Adobe	2017, 2018
	University of Michigan Rackham Merit Fellowship	2017
	ICLR, Neurips BAI Travel Award	2017, 2019
	NSF Graduate Research Fellowship (Neuroscience)	2015
	Provost Award for Academic Excellence (20/5000 graduates chosen)	2015
	Researcher of the Month (1 monthly at Stony Brook University)	2014
	HHMI Minority Undergraduate Research Fellowship	2014
	$\Sigma\Pi\Sigma$ Physics Honor Society (sponsored by Alfred Goldhaber)	2013
	NSF Louis Stokes Alliance for Minority Participation Scholar	2011

PREPRINTS

Wilka Carvalho, Samuel J. Gershman. “*Preemptive Solving of Future Problems: Multitask Preplay in Humans and Machines.*” 2025

Wilka Carvalho*, Andrew Lampinen*. “*Naturalistic Computational Cognitive Science.*” 2025

Wilka Carvalho. “*NiceWebRL: a Python library for human subject experiments with reinforcement learning environments.*” 2025

Aneesh Muppidi, Samuel J. Gershman, **Wilka Carvalho**. “*Unsupervised Agent-centric Representation Learning with Variational Agent Inference.*” 2025

PUBLICATIONS

Kunal Jha, **Wilka Carvalho**, Natasha Jaques, Max Kleiman-Weiner. “*Cross-environment Cooperation Enables Zero-shot Multi-agent Coordination.*” In ICML, 2025 (**Oral**)

Wilka Carvalho*, Momchil S Tomov*, William de Cothi*, Caswell Barry, Samuel J. Gershman. “*Predictive representations: building blocks of intelligence.*” In Neural Computation, 2024

Wilka Carvalho, Andre Saraiva, Angelos Filos, Andrew Lampinen, Loic Matthey, Richard L Lewis, Honglak Lee, Satinder Singh, Danilo Jimenez Rezende, Daniel Zoran. “*Combining behaviors with the successor features keyboard.*” In NeurIPS, 2024

Wilka Carvalho, Andrew Lampinen, Kyriacos Nikiiforou, Felix Hill, Murray Shanahan. “*Feature-Attending Recurrent Modules for Generalization in Reinforcement Learning.*” In TMLR, 2023

Wilka Carvalho, Angelos Filos, Richard L. Lewis, Honglak Lee, Satinder Singh. “*Composing Task Knowledge with Modular Successor Feature Approximators.*” In ICLR, 2023

Lajanugen Logeswaran, **Wilka Carvalho**, Honglak Lee. “*Learning Compositional Tasks from Language Instructions.*” In AACL, 2023

Chris Hoang, Sungryull Sohn, Jongwook Choi, **Wilka Carvalho**, Honglak Lee. “*Successor Feature Landmarks for Long-Horizon Goal-Conditioned Reinforcement Learning.*” In NeurIPS, 2021

Wilka Carvalho, Anthony Liang, Kimin Lee, Sungryull Sohn, Honglak Lee, Richard L. Lewis, Satinder Singh. “*Reinforcement Learning for Sparse-Reward Object-Interaction Tasks in First-person Simulated 3D Environments.*” In IJCAI, 2021

Wilka Carvalho*, Sanjay Purushotham*, Tanachat Nilanon, Yan Liu. “*Variational Recurrent Adversarial Domain Adaptation.*” In ICLR, 2017

WORKSHOP PUBLICATIONS

Lajanugen Logeswaran, **Wilka Carvalho**, Honglak Lee. “*Learning Compositional Tasks from Language Instructions.*” In NeurIPS Deep RL Workshop, 2021

Wilka Carvalho, Murray Shanahan. “*Learning to Represent State with Perceptual Schemata.*” In ICML {Unsupervised RL, Real Life RL} Workshops, 2021

Wilka Carvalho, Anthony Liang, Kimin Lee, Sungryull Sohn, Honglak Lee, Richard L. Lewis, Satinder Singh. “*Reinforcement Learning for Sparse-Reward Object-Interaction Tasks in First-person Simulated 3D Environments.*” In {ICML Object-Oriented Learning, NeurIPS DeepRL} Workshops, 2020

Chris Hoang, Sungryull Sohn, Jongwook Choi, **Wilka Carvalho**, Honglak Lee. “*Successor Landmarks for Efficient Exploration and Long-Horizon Navigation.*” In NeurIPS DeepRL Workshop,

2020

Wilka Carvalho, Kimin Lee, Anthony Liang, Ryan Krueger, Richard L. Lewis, Satinder Singh, Honglak Lee. “Efficiently Learning to Perform Household Task with Object-oriented Exploration.” In NeurIPS BAI, 2019 (**Oral**)

Bryant Chen*, **Wilka Carvalho***, Benjamin Edwards, Taesung Lee, Ian Molloy, Heiko Ludwig. “Detecting Backdoor Attacks on Deep Neural Networks by Activation Clustering.” In AAAI Artificial Intelligence Safety Workshop, 2018 (**Best Paper**)

Sanjay Purushotham*, **Wilka Carvalho***, Yan Liu. “Variational Adversarial Deep Domain Adaptation for Health Care Time Series Analysis.” In NeurIPS Machine Learning for Healthcare Workshop, 2016 (**Spotlight**)

Wilka Carvalho. “Modeling a Detection of internally reflected Cherenkov light (DIRC) Particle Detector for High-Multiplicity Collisions.” In SURC, 2015

TALKS

NiceWebRL: A Python Library for Human Experiments with Machine RL Environments

Human-AI Complementarity for Decision Making Workshop. *Carnegie Mellon University*. (September, 2025)

Cognitive Computational Neuroscience Conference. *Amsterdam*. (August, 2025)

Naturalistic Computational Cognitive Science

Cognitive Computational Neuroscience Conference. *Amsterdam*. (August, 2025)

Scaling Up Reinforcement Learning Theories of Cognition

Cognition, Brain, and Behavior Seminar. *Harvard University*. (April, 2025)

Center for Theoretical Neuroscience. *Columbia University*. (April, 2025)

Mattar Lab. *NYU*. (April, 2025)

Emerge Lab. *NYU*. (April, 2025)

Computation and Decision-Making Lab. *NYU*. (April, 2025)

RL Seminar. *Tübingen*. (March, 2025)

Google DeepMind. *San Francisco*. (March, 2025)

Cognition Seminar. *UC Berkeley*. (March, 2025)

The Rational Altruism Lab. *UCLA*. (March, 2025)

Social and Decision Neuroscience Seminar. *Caltech*. (March, 2025)

Symposium on “Frontiers of Machine Learning and AI. *USC*. (March, 2025)

Neurolunch Seminar. *Harvard*. (March, 2025)

RL Seminar. *UMich*. (February, 2025)

Vision Lab. *Harvard University*. (February, 2025)

Sabatini Lab. *Harvard University*. (February, 2025)

Predictive Representations: Building Blocks of Intelligence

Computational Cognitive Science Laboratory. *UC Berkeley*. (March, 2025)

Computational Cognitive Neuroscience Lab. *Harvard University*. (February, 2025)

Neuro Systems Club. *Harvard University*. (Fall, 2024)

Mathis Group for Computational Neuroscience and AI. *EPFL*. (Fall, 2024)

Theory of Learning Lab. *Gatsby Computational Neuroscience Unit (UCL)*. (Summer, 2024)

Conference on the Mathematics of Neuroscience and AI. *Rome, Italy*. (Summer, 2024)

Theoretical Cognitive Neuroscience Lab. *Boston University*. (Fall, 2023)

Towards Human-like Deep Reinforcement Learning

Computational Cognitive Science Group. *Princeton*. (Fall, 2023)

Kempner Institute. *Harvard University*. (March, 2023)

Composing Task Knowledge with Modular Successor Feature Approximators
RL at Harvard. *Harvard University*. (August, 2023)

Weak Inductive Biases for Composable Primitive Representations
Computational Neuroscience Group. *UCL*. (October, 2022)
Intelligent Robot Lab. *Brown*. (February, 2022)
Computational Cognitive Neuroscience Lab. *Harvard University*. (February, 2022)
Computational Cognitive Neuroscience Lab. *University of California—Berkeley*. (January, 2022)
Deep Learning Community of Practice. *Intel*. (January, 2022)
Deep RL Workshop. *NeurIPS*. (December, 2021)

Reinforcement Learning for Sparse-Reward Object-Interaction Tasks in a First-Person Simulated 3D Environment

Object Representations for Learning Workshop. *NeurIPS*. (December, 2020)
Cognitive Tools Lab. *University of California—San Diego*. (November, 2020)
Cognitive Science Community. *University of Michigan*. (November, 2020)
Stony Brook Society of Physics. *Stony Brook, NY*. (April, 2020)

Variational Recurrent Adversarial Deep Domain Adaptation

Machine Learning Lunch Seminar. *University of Southern California*. (April, 2017)

PATENTS

Wilka Carvalho, Angelos Filos, Danilo Jimenez Rezende, Daniel Zoran. “*Controlling Agents by Transferring Successor Features To New tasks*.” 2023

Wilka Carvalho. “*Compositional Generalization in Reinforcement Learning*.” 2023

Bryant Chen, **Wilka Carvalho**, Heiko Ludig, Ian Molloy, Jialong Zhang, Benjamin Edwards. “*Detecting poisoning attacks on neural networks by activation clustering*.” 2020

Wilka Carvalho, Bryant Chen, Benjamin Edwards, Taesung Lee, Ian Molloy, Jialong Zhang. “*Using Gradients to Detect Backdoors in Neural Networks*.” 2018

Wilka Carvalho, Yan Liu, Tanachat Nilanon, Sanjay Purushotham. “*Effective Knowledge Transfer Among Patient Populations via Deep Learning*.” 2017

RESEARCH EXPERIENCE

University of Michigan—Ann Arbor, Ann Arbor, Michigan USA

AI Lab, August 2018 – Present

Advisors: [Honglak Lee](#), [Satinder Singh](#), [Richard Lewis](#)

- Studying how deep reinforcement learning agents can discover representations and temporally-extended behaviors that enable generalization

DeepMind, London, UK

Generative Models Team, September 2022 – May 2023

Advisor: [Danilo Rezende](#)

- Studying transfer in 3D environments with successor features

DeepMind, London, UK

Cognition Team, March 2021 – July 2021

Advisor: [Murray Shanahan](#)

- Studying agent-state representations for generalization of deep reinforcement learning agents

in 3D environments

Microsoft Research, Redmond, Washington USA

Medical Devices Group, *June 2018 – August 2018*

Advisor: [Sumit Basu](#)

- Developed a machine learning model that can predict clinical measures from physiological signals measured by a wearable device.

IBM Research, San Jose, California USA

AI Platform Research Group, *September 2017 – December 2017*

Advisor: [Heiko Ludwig](#)

- Contributed to novel research algorithm by suggesting subspace projection technique that increased our performance from 15% to 95% accuracy.
- Research led to workshop publication and was presented in global meeting to IBM executives including CEO and selected for funding.

Visa Research, Palo Alto, California USA

Data Analytics Group, *June 2017 – August 2017*

Advisor: [Hao Yang](#)

- Implemented model and baselines for language generation and question answering.

University of Southern California, Los Angeles, California USA

Melady Machine Learning Lab, *November 2015 – May 2017*

Advisor: [Yan Liu](#)

Samsung and NSF funded project: “Variational Adversarial Deep Domain Adaptation for Health Care Time Series Analysis”

- Implemented novel neural network that employed variational inference and adversarial training for transfer learning of multivariate time-series.
- Research led to a publication and a patent.
- Communicated research to general public through research feature by the USC Graduate School and to technical audience at ICLR poster presentation.

Stony Brook University, Stony Brook, New York USA

Heavy Ion Research Group, *January 2013 – August 2015*

Advisor: [Axel Drees](#)

DOE funded project: “Modeling a Detection of internally reflected Cherenkov light Particle Detector for High-Multiplicity Collisions”

- Contributed methods from multivariate calculus and linear algebra to particle detection algorithm. Accuracy improved from 60% to 80%.
- Designed and implemented a statistical analysis pipeline in C++ for measuring efficacy of particle detection algorithm.

Stony Brook University, Stony Brook, New York USA

Computational Neuroscience Group, *Fall 2014*

Advisor: [Giancarlo La Camera](#)

NSF LSAMP funded project: “Spectral Analysis of Rodent Neural Data”

- Performed spectral analyses on neural data to determine behavioral correlates of neural activity.

California Institute of Technology, Pasadena, California USA

Emotion and Social Cognition Laboratory, *Summer 2014*

Advisor: [Ralph Adolphs](#)

HHMI funded project: “Modeling the Cognitive Process of Attributing Traits to Others”

- Formulated a trait learning behavioral experiment to study human inference.

University of Minnesota, Minneapolis, Minnesota USA

Neuromodulation Research and Technology Laboratory, *Summer 2013*

Advisor: [Matthew Johnson](#)

NIH funded project: “Modeling Deep Brain Stimulation of Globus Palidus Internus”

- Implemented python script to build a biologically feasible computational model of neural networks

National Central University, Jhongli City, Taiwan

Turbulent Combustion Laboratory, *Summer 2012*

Advisor: [Shenqyang Shy](#)

TEACHING EXPERIENCE

Stony Brook University, Stony Brook, NY

Calculus Instructor, *Spring 2015*

Worked with two math professors to develop and teach a supplementary calculus curriculum that promoted minority representation in stem majors.

Stony Brook University, Stony Brook, NY

Educational Opportunity Program Personal Tutor, *Spring 2013 - Fall 2014*

Tutored marginalized students in introductory physics and math courses

SERVICE

Action Editor, Transactions on Machine Learning Research, 2023

Area Chair, Reinforcement Learning Conference, 2023

Reviewer, Neural Computation, 2023

Reviewer, NeurIPS, 2020, 2021, 2022, 2023

Reviewer, ICLR, 2021, 2022

Reviewer, ICML, 2021, 2022

Student Volunteer, ICLR, 2017

OUTREACH

HS2 Engineer's Highschool Panel, Virtual, 2021

Stony Brook Society of Physics, Stony Brook, NY, 2020

Michigan Explore Graduate Studies Workshop, Ann Arbor, MI, 2019

Research and Fellowships Week NSF Panel, Los Angeles, CA, 2016

National Society of Black Engineers Grad Panel, Los Angeles, CA, 2016

Graduate School External Fellowship Boot Camp, Los Angeles, CA, 2016

Mentored marginalized high school youth through the Pullias Center for Higher Education, Los Angeles, CA, 2016

Engineering Graduate Diversity Symposium, Los Angeles, CA, 2015

Black Student Association: What it takes to go to Graduate School, Los Angeles, CA, 2015

Collegiate Science and Technology Entry Program Undergraduate Research Panel, Stony Brook, CA, 2014

SKILLS

Machine Learning Software: Pytorch, TensorFlow, Theano, Keras

Neuroscience Software: Neuron

Languages: Python, C++, C, Java

Systems: Unix, Linux, OSX

PRESS

[Finding Gaps in Your Grad School Apps](#)

[This Week In Machine Learning \(TWIML\) Research Feature](#)

University of Michigan Research Feature
Exploring the source of social stereotypes
Black History Month: Why a career in science?
Research Feature by the USC Graduate School
2015 NSF Graduate Research Fellow Wilka Carvalho
Biomath Learning Center Launches Modified Supplemental Instruction Program
URECA Research of the Month: Wilka Carvalho
Student Feature by Stony Brook University

INTERESTS

• traveling • chess • software development • improvisational dance • deadpan humor